


## 4.1


# Output and Loss Functions

# Output and Loss Functions



- ▶ The **output layer** is the last layer in a neural network which computes the output
- ▶ The **loss function** compares the result of the output layer to the target value(s)
- ▶ Choice of output layer and loss function depends on task (discrete, continuous, ..)

# Output and Loss Functions




- ▶ The **output layer** is the last layer in a neural network which computes the output
- ▶ The **loss function** compares the result of the output layer to the target value(s)
- ▶ Choice of output layer and loss function depends on task (discrete, continuous, ..)

# Loss Function

What is the goal of optimizing the loss function?



- Tries to make the **model output** (=prediction) similar to the **target** (=data)
- Think of the loss function as a **measure of cost** being paid for a prediction



# Loss Function

What is the goal of optimizing the loss function?

- ▶ Tries to make the **model output** (=prediction) similar to the **target** (=data)
- ▶ Think of the loss function as a **measure of cost** being paid for a prediction




# Loss Function

How to design a good loss function?

- ▶ A loss function can be any differentiable function that we wish to optimize
- ▶ Deriving the cost function from the **maximum likelihood principle** removes the burden of manually designing the cost function for each model
- ▶ Consider the output of the neural network as **parameters of a distribution** over  $y_i$

$$\begin{aligned}\hat{\mathbf{w}}_{ML} &= \underset{\mathbf{w}}{\operatorname{argmax}} p_{model}(\mathbf{y}|\mathbf{X}, \mathbf{w}) \\ &\stackrel{\text{iid}}{=} \underset{\mathbf{w}}{\operatorname{argmax}} \prod_{i=1}^N p_{model}(y_i|\mathbf{x}_i, \mathbf{w}) \\ &= \underset{\mathbf{w}}{\operatorname{argmax}} \underbrace{\sum_{i=1}^N \log p_{model}(y_i|\mathbf{x}_i, \mathbf{w})}_{\text{Log-Likelihood}}\end{aligned}$$

# Loss Function



## Example:

- Neural network  $f_{\mathbf{w}}(\mathbf{x})$  predicts mean  $\mu$  of Gaussian distribution over  $y$ :

$$p(y|\mathbf{x}, \mathbf{w}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - f_{\mathbf{w}}(\mathbf{x}))^2}{2\sigma^2}\right)$$


- We want to maximize the probability of the target  $y$  under this distribution

# Recap: Regression

Input



Model




Output

143,52 €

► Mapping:  $f_{\mathbf{w}} : \mathbb{R}^N \rightarrow \mathbb{R}$




## Recap: Binary Classification



► **Mapping:**  $f_{\mathbf{w}} : \mathbb{R}^{W \times H} \rightarrow \{\text{"Beach"}, \text{"No Beach"}\}$

## Recap: Multi-Class Classification



► **Mapping:**  $f_{\mathbf{w}} : \mathbb{R}^{W \times H} \rightarrow \{\text{"Beach"}, \text{"Mountain"}, \text{"City"}, \text{"Forest"}\}$


# Regression Problems

# Gaussian Distribution

## Gaussian distribution:

$$p(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$

- ▶  $\mu$  : mean
- ▶  $\sigma$  : standard deviation
- ▶ The distribution has thin “tails”:  
 $p(y) \rightarrow 0$  quickly as  $y \rightarrow \infty$
- ▶ It thus penalizes outliers strongly



## Gaussian Distribution / $L_2$ Loss

Let  $p_{model}(y|\mathbf{x}, \mathbf{w}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y-f_{\mathbf{w}}(\mathbf{x}))^2}{2\sigma^2}\right)$  be a **Gaussian distribution**. We obtain:

$$\begin{aligned}\hat{\mathbf{w}}_{ML} &= \operatorname{argmax}_{\mathbf{w}} \sum_{i=1}^N \log p_{model}(y_i|\mathbf{x}_i, \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{w}} - \sum_{i=1}^N \frac{1}{2} \log(2\pi\sigma^2) - \sum_{i=1}^N \frac{1}{2\sigma^2} (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 \\ &= \operatorname{argmax}_{\mathbf{w}} - \sum_{i=1}^N (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 \\ &= \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^N \underbrace{(f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2}_{L_2 \text{ Loss}}\end{aligned}$$


In other words, we minimize the **squared loss** ( $=L_2$  loss), affected strongly by outliers.

# Laplace Distribution

## Laplace distribution:

$$p(y) = \frac{1}{2b} \exp\left(-\frac{|y - \mu|}{b}\right)$$

- ▶  $\mu$  : location
- ▶  $b$  : scale
- ▶ The distribution has heavy “tails”:  
 $p(y) \rightarrow 0$  more slowly as  $y \rightarrow \infty$
- ▶ Penalizes outliers less strongly
- ▶ Thus often preferred in practice



## Laplace Distribution / $L_1$ Loss

Let  $p_{model}(y|\mathbf{x}, \mathbf{w}) = \frac{1}{2b} \exp\left(-\frac{|y-f_{\mathbf{w}}(\mathbf{x})|}{b}\right)$  be a **Laplace distribution**. We obtain:

$$\begin{aligned}\hat{\mathbf{w}}_{ML} &= \operatorname{argmax}_{\mathbf{w}} \sum_{i=1}^N \log p_{model}(y_i|\mathbf{x}_i, \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{w}} - \sum_{i=1}^N \log(2b) - \sum_{i=1}^N \frac{1}{b} |f_{\mathbf{w}}(\mathbf{x}_i) - y_i| \\ &= \operatorname{argmax}_{\mathbf{w}} - \sum_{i=1}^N |f_{\mathbf{w}}(\mathbf{x}_i) - y_i| \\ &= \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^N \underbrace{|f_{\mathbf{w}}(\mathbf{x}_i) - y_i|}_{L_1 \text{ Loss}}\end{aligned}$$

We minimize the **absolute loss** ( $=L_1$  loss) which is more robust than  $L_2$ .

# Predicting all Parameters

Let  $p_{model}(y|\mathbf{x}, \mathbf{w}) = \frac{1}{2 g_{\mathbf{w}}(\mathbf{x})} \exp\left(-\frac{|y - f_{\mathbf{w}}(\mathbf{x})|}{g_{\mathbf{w}}(\mathbf{x})}\right)$  be a **Laplace distribution**. We obtain:

$$\begin{aligned}\hat{\mathbf{w}}_{ML} &= \operatorname{argmax}_{\mathbf{w}} \sum_{i=1}^N \log p_{model}(y_i|\mathbf{x}_i, \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{w}} - \sum_{i=1}^N \log(2 g_{\mathbf{w}}(\mathbf{x}_i)) - \sum_{i=1}^N \frac{1}{g_{\mathbf{w}}(\mathbf{x}_i)} |f_{\mathbf{w}}(\mathbf{x}_i) - y_i|\end{aligned}$$


In this case, we predict both the **location**  $\mu$  and the **scale**  $b$  with a neural network.  $f_{\mathbf{w}}(\mathbf{x}_i)$  and  $g_{\mathbf{w}}(\mathbf{x})$  are typically the same except for the last layer. Allows for estimating the **aleatoric uncertainty** (observation noise) with the neural network itself.




# Predicting all Parameters




Input Image






Ground Truth



Segmentation



Aleatoric Uncertainty




# Mixture Density Networks

To represent multi-modal distributions, we can also model **mixture densities**:


$$p_{model}(y|\mathbf{x}, \mathbf{w}) = \sum_{m=1}^M \pi_m \frac{1}{2 g_{\mathbf{w}}^{(m)}(\mathbf{x})} \exp \left( -\frac{|y - f_{\mathbf{w}}^{(m)}(\mathbf{x})|}{g_{\mathbf{w}}^{(m)}(\mathbf{x})} \right)$$

## Example:

- ▶ Mixture of Laplace distribution
- ▶  $\pi_m \in [0, 1]$ : weight of mode  $m$
- ▶ Constraint  $\sum_m \pi_m = 1$
- ▶ Location  $\mu_m$  and scale  $b_m$  of each mode  $m$  modeled by neural network




# Output Layer for Regression Problems



- For most outputs (e.g.,  $\mu \in \mathbb{R}$ ), the output layer is linear (no non-linearity)
- For some outputs (e.g.,  $b \in \mathbb{R}^+$ ), we need a squashing function (ReLU, softplus)

# Loss Function for Regression Problems



- ▶ Gaussian/Laplacian model distribution correspond to  $L_2$  and  $L_1$  loss functions
- ▶ It is also possible to predict uncertainty (variance/scale) or multiple modes (MDN)

# Classification Problems


# Image Classification



## **MNIST Handwritten Digits:**

- ▶ One of the most popular datasets in ML (many variants, still in use today)
- ▶ Based on a data from the National Institute of Standards and Technology
- ▶ Hand written by Census Bureau employees and high-school children
- ▶ Resolution: 28 x 28 pixels, 60k training samples with labels, 10k test samples

# Image Classification



## Curse of Dimensionality:


- ▶ There exist  $2^{784} = 10^{236}$  possible binary images of resolution 28 x 28 pixels
- ▶ MNIST is gray-scale, thus  $256^{784}$  combinations  $\Rightarrow$  impossible to enumerate
- ▶ Why is image classification with just 60k labeled training images even possible?
- ▶ Answer: Images concentrated on low-dimensional manifold in  $\{1, \dots, 256\}^{784}$

# Bernoulli Distribution

## Bernoulli distribution:

$$p(y) = \mu^y (1 - \mu)^{(1-y)}$$

- ▶  $\mu$ : probability for  $y = 1$
- ▶ Handles only two classes  
e.g. ("cats" vs. "dogs")





## Bernoulli Distribution / BCE Loss

Let  $p_{model}(y|\mathbf{x}, \mathbf{w}) = f_{\mathbf{w}}(\mathbf{x})^y (1 - f_{\mathbf{w}}(\mathbf{x}))^{(1-y)}$  be a **Bernoulli distribution**. We obtain:

$$\begin{aligned}\hat{\mathbf{w}}_{ML} &= \operatorname{argmax}_{\mathbf{w}} \sum_{i=1}^N \log p_{model}(y_i|\mathbf{x}_i, \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{w}} \sum_{i=1}^N \log \left[ f_{\mathbf{w}}(\mathbf{x}_i)^{y_i} (1 - f_{\mathbf{w}}(\mathbf{x}_i))^{(1-y_i)} \right] \\ &= \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^N \underbrace{-y_i \log f_{\mathbf{w}}(\mathbf{x}_i) - (1 - y_i) \log(1 - f_{\mathbf{w}}(\mathbf{x}_i))}_{\text{BCE Loss}}\end{aligned}$$

In other words, we minimize the **binary cross-entropy (BCE)** loss.

Remark: Last layer of  $f_{\mathbf{w}}(\mathbf{x})$  can be a sigmoid function such that  $f_{\mathbf{w}}(\mathbf{x})^y \in [0, 1]$ .

How can we scale this to multiple classes?

# Categorical Distribution

## Categorical distribution:


$$p(y = c) = \mu_c$$

- ▶  $\mu_c$ : probability for class  $c$
- ▶ Multiple classes, multiple modes





## Alternative notation:

$$p(\mathbf{y}) = \prod_{c=1}^C \mu_c^{y_c}$$

- ▶  $\mathbf{y}$ : “one-hot” vector with  $y_c \in \{0, 1\}$
- ▶  $\mathbf{y} = (0, \dots, 0, 1, 0, \dots, 0)^\top$  with all zeros except for one (the true class)



# One-Hot Vector Representation

class	$y$	$\mathbf{y}$
	1	$(1, 0, 0, 0)^\top$
	2	$(0, 1, 0, 0)^\top$
	3	$(0, 0, 1, 0)^\top$
	4	$(0, 0, 0, 1)^\top$

- ▶ One-hot vector  $\mathbf{y}$  with binary elements  $y_c \in \{0, 1\}$
- ▶ Index  $c$  with  $y_c = 1$  determines the correct class, and  $y_k = 0$  for  $k \neq c$
- ▶ Interpretation as discrete distribution with all probability mass at the true class
- ▶ Often used in ML as it can make formalism more convenient

# Categorical Distribution / CE Loss

Let  $p_{model}(\mathbf{y}|\mathbf{x}, \mathbf{w}) = \prod_{c=1}^C f_{\mathbf{w}}^{(c)}(\mathbf{x})^{y_c}$  be a **Categorical distribution**. We obtain:

$$\begin{aligned}\hat{\mathbf{w}}_{ML} &= \underset{\mathbf{w}}{\operatorname{argmax}} \sum_{i=1}^N \log p_{model}(\mathbf{y}_i|\mathbf{x}_i, \mathbf{w}) \\ &= \underset{\mathbf{w}}{\operatorname{argmax}} \sum_{i=1}^N \log \prod_{c=1}^C f_{\mathbf{w}}^{(c)}(\mathbf{x}_i)^{y_{i,c}} \\ &= \underset{\mathbf{w}}{\operatorname{argmin}} \underbrace{\sum_{i=1}^N \sum_{c=1}^C -y_{i,c} \log f_{\mathbf{w}}^{(c)}(\mathbf{x}_i)}_{\text{CE Loss}}\end{aligned}$$

In other words, we minimize the **cross-entropy (CE)** loss.

The target  $\mathbf{y} = (0, \dots, 0, 1, 0, \dots, 0)^\top$  is a “one-hot” vector with  $y_c$  its  $c$ ’th element.

# Softmax

How can we ensure that  $f_{\mathbf{w}}^{(c)}(\mathbf{x})$  predicts a **valid Categorical (discrete) distribution**?

- ▶ We must guarantee (1)  $f_{\mathbf{w}}^{(c)}(\mathbf{x}) \in [0, 1]$  and (2)  $\sum_{c=1}^C f_{\mathbf{w}}^{(c)}(\mathbf{x}) = 1$
- ▶ An element-wise sigmoid as output function would ensure (1) but not (2)
- ▶ Solution: The **softmax function** guarantees both (1) and (2):

$$\text{softmax}(\mathbf{x}) = \left( \frac{\exp(x_1)}{\sum_{k=1}^C \exp(x_k)}, \dots, \frac{\exp(x_C)}{\sum_{k=1}^C \exp(x_k)} \right)$$

- ▶ Let  $\mathbf{s}$  denote the network output after the last affine layer (=scores). Then:

$$f_{\mathbf{w}}^{(c)}(\mathbf{x}) = \frac{\exp(s_c)}{\sum_{k=1}^C \exp(s_k)} \quad \Rightarrow \quad \log f_{\mathbf{w}}^{(c)}(\mathbf{x}) = s_c - \log \sum_{k=1}^C \exp(s_k)$$

- ▶ Remark:  $s_c$  is a direct contribution to the loss function, i.e., it does not saturate



# Log Softmax

**Intuition:** Assume  $c$  is the correct class. Our goal is to maximize the log softmax:

$$\log f_{\mathbf{w}}^{(c)}(\mathbf{x}) = s_c - \log \sum_{k=1}^C \exp(s_k)$$

- ▶ The first term encourages the score  $s_c$  for the correct class  $c$  to increase
- ▶ The second term encourages all scores in  $\mathbf{s}$  to decrease
- ▶ The second term can be approximated by:  $\log \sum_{k=1}^C \exp(s_k) \approx \max_k s_k$   
as  $\exp(s_k)$  is insignificant for all  $s_k < \max_k s_k$
- ▶ Thus, the loss always strongly penalizes the most active incorrect prediction
- ▶ If the correct class already has the largest score (i.e.,  $s_c = \max_k s_k$ ), both terms roughly cancel and the example will contribute little to the overall training cost

# Log Softmax Example



- The second term becomes:  $\log \sum_{k=1}^C \exp(s_k) = 4.06 \approx s_3 = \max_k s_k$
- For  $c = 2$  we obtain:  $\log f_{\mathbf{w}}^{(c)}(\mathbf{x}) = s_c - \log \sum_{k=1}^C \exp(s_k) = 1 - 4.06 \approx -3$
- For  $c = 3$  we obtain:  $\log f_{\mathbf{w}}^{(c)}(\mathbf{x}) = s_c - \log \sum_{k=1}^C \exp(s_k) = 4 - 4.06 \approx 0$



# Softmax

- ▶ Predicting  $C$  values/scores overparameterizes the Categorical distribution
- ▶ As the distribution sums to 1 only  $C-1$  parameters are necessary
- ▶ Example: Consider  $C = 2$  and fix one degree of freedom ( $x_2 = 0$ ):





$$\begin{aligned}\text{softmax}(\mathbf{x}) &= \left( \frac{\exp(x_1)}{\exp(x_1) + \exp(x_2)}, \frac{\exp(x_2)}{\exp(x_1) + \exp(x_2)} \right) \\ &= \left( \frac{\exp(x_1)}{\exp(x_1) + 1}, \frac{1}{\exp(x_1) + 1} \right) \\ &= \left( \frac{1}{1 + \exp(-x_1)}, 1 - \frac{1}{1 + \exp(-x_1)} \right) \\ &= (\sigma(x_1), 1 - \sigma(x_1))\end{aligned}$$

- ▶ The softmax is a **multi-class generalization** of the sigmoid function
- ▶ In practice, the overparameterized version is often used (simpler to implement)

# Softmax

$$\text{softmax}(\mathbf{s}) = \left( \frac{\exp(s_1)}{\sum_{k=1}^C \exp(s_k)}, \dots, \frac{\exp(s_C)}{\sum_{k=1}^C \exp(s_k)} \right)$$

- ▶ The name “softmax” is confusing, “soft argmax” would be more precise as it is a continuous and differentiable version of argmax (in one-hot representation)
- ▶ Example with four classes:



# Softmax

- ▶ Softmax responds to differences between inputs
- ▶ It is invariant to adding the same scalar to all its inputs:

$$\text{softmax}(\mathbf{x}) = \text{softmax}(\mathbf{x} + c)$$

- ▶ We can therefore derive a numerically more stable variant:

$$\text{softmax}(\mathbf{x}) = \text{softmax}(\mathbf{x} - \max_{k=1..L} x_k)$$





- ▶ Allows accurate computation even when  $\mathbf{x}$  is large
- ▶ Illustrates again that softmax depends on differences between scores

# Cross Entropy Loss with Softmax Example

**Putting it together:** Cross Entropy Loss for a single training sample  $(\mathbf{x}, \mathbf{y}) \in \mathcal{X}$ :


$$\text{CE Loss: } \sum_{c=1}^C -y_c \log f_{\mathbf{w}}^{(c)}(\mathbf{x})$$

Example: Suppose  $C = 4$  and 4 training samples  $\mathbf{x}$  with labels  $\mathbf{y}$

Input $\mathbf{x}$	Label $\mathbf{y}$	Predicted scores $\mathbf{s}$	softmax( $\mathbf{s}$ )	CE Loss
	$(\mathbf{1}, 0, 0, 0)^T$	$(+3, +1, -1, -1)^T$	$(\mathbf{0.85}, 0.12, 0.02, 0.02)^T$	0.16
	$(0, \mathbf{1}, 0, 0)^T$	$(+3, +3, +1, +0)^T$	$(0.46, \mathbf{0.46}, 0.06, 0.02)^T$	0.78
	$(0, 0, \mathbf{1}, 0)^T$	$(+1, +1, +1, +1)^T$	$(0.25, 0.25, \mathbf{0.25}, 0.25)^T$	1.38
	$(0, 0, 0, \mathbf{1})^T$	$(+3, +2, +3, -1)^T$	$(0.42, 0.16, 0.42, \mathbf{0.01})^T$	4.87


- Sample 4 contributes most strongly to the loss function! (elephant in the room)

# Output Layer for Classification Problems



- For 2 classes, we can predict 1 value and use a sigmoid, or 2 values with softmax
- For  $C > 2$  classes we typically predict  $C$  scores and use a softmax non-linearity


# Loss Function for Classification Problems



- For 2 classes, we use the binary cross-entropy loss (BCE)
- For  $C > 2$  classes, we use the cross-entropy loss (CE)



# Activation Functions



- ▶ Hidden layer  $\mathbf{h}_i = g(\mathbf{A}_i \mathbf{h}_{i-1} + \mathbf{b}_i)$  with **activation function**  $g(\cdot)$  and weights  $\mathbf{A}_i, \mathbf{b}_i$
- ▶ The activation function is frequently applied **element-wise** to its input
- ▶ Activation functions must be **non-linear** to learn non-linear mappings
- ▶ Some of them are not differentiable everywhere (but still ok for training)



# Activation Functions


## Sigmoid:

$$g(x) = \frac{1}{1 + \exp(-x)}$$

- ▶ Maps input to range  $[0, 1]$
- ▶ Neuroscience interpretation as saturating “firing rate” of neurons


## Problems:

- ▶ Saturation “kills” gradients
- ▶ Outputs are not zero-centered
- ▶ Introduces bias after first layer




# Activation Functions

## Sigmoid Problem #1:



- ▶ Downstream gradient becomes zero when input  $x$  is saturated:  $g'(x) \approx 0$
- ▶ No learning if  $x$  is very small ( $< -10$ )
- ▶ No learning if  $x$  is very large ( $> 10$ )



# Activation Functions

## Sigmoid Problem #2:


$$g(x) = \frac{1}{1 + \exp(-x)} \quad x = \sum_i a_i x_i + b$$

- ▶ Sigmoid is always positive  $\Rightarrow x_i$  also
- ▶ Gradient of sigmoid is always positive

The gradient wrt. parameter  $a_i$  is given by:



$$\frac{\partial \mathcal{L}}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} \frac{\partial x}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} x_i$$

- ▶ Therefore:  $\text{sgn}(\partial \mathcal{L} / \partial a_i) = \text{sgn}(\partial \mathcal{L} / \partial g)$
- ▶ All gradients have the same sign (+ or -)



# Activation Functions

## Sigmoid Problem #2:



- Restricts gradient updates and leads to inefficient optimization (minibatches help)

# Activation Functions


## Tanh:

$$g(x) = \frac{2}{1 + \exp(-2x)} - 1$$

- ▶ Maps input to range  $[-1, 1]$
- ▶ Anti-symmetric
- ▶ Zero-centered

## Problems:

- ▶ Again, saturation “kills” gradients



# Activation Functions


## Rectified Linear Unit (ReLU):

$$g(x) = \max(0, x)$$

- ▶ Does not saturate (for  $x > 0$ )
- ▶ Leads to fast convergence
- ▶ Computationally efficient


## Problems:

- ▶ Not zero-centered
- ▶ No learning for  $x < 0 \Rightarrow$  dead ReLUs




# Activation Functions

## ReLU Problem:



- ▶ Downstream gradient becomes zero when input  $x < 0$
- ▶ Results in so-called “dead ReLUs” that never participate in learning
- ▶ Often initialize with pos. bias ( $b > 0$ )




# Activation Functions

## Leaky ReLU:

$$g(x) = \max(0.01x, x)$$

- ▶ Does not saturate (i.e., will not die)
- ▶ Closer to zero-centered outputs
- ▶ Leads to fast convergence
- ▶ Computationally efficient






# Activation Functions

## Parametric ReLU:

$$g(x) = \max(\alpha x, x)$$

- ▶ Does not saturate (i.e., will not die)
- ▶ Leads to fast convergence
- ▶ Computationally efficient
- ▶ Parameter  $\alpha$  learned from data




# Activation Functions




## Exponential Linear Units (ELU):

$$g(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$

- ▶ All benefits of Leaky ReLU
- ▶ Adds some robustness to noise
- ▶ Default  $\alpha = 1$



# Activation Functions



**Maxout:**  $g(x) = \max(\mathbf{a}_1^\top \mathbf{x} + b_1, \mathbf{a}_2^\top \mathbf{x} + b_2)$

- Generalizes ReLU and Leaky ReLU
- Increases the number of parameters per neuron

# Activation Functions

## **Summary:**

- ▶ No one-size-fits-all: Choice of activation function depends on problem
- ▶ We only showed the most common ones, there exist many more
- ▶ Best activation function/model is often found using trial-and-error in practice
- ▶ It is important to ensure a good “gradient flow” during optimization

## **Rule of Thumb:**

- ▶ Use ReLU by default (with small enough learning rate)
- ▶ Try Leaky ReLU, Maxout, ELU for some small additional gain
- ▶ Prefer Tanh over Sigmoid (Tanh often used in recurrent models)

# Implementation


# Numerical Differentiation

- ▶ Murphy: “Anything that can go wrong will.”
- ▶ When implementing the backward pass of activation, output or loss functions it is important to ensure that all gradients are correct!
- ▶ Verify via Newton’s difference quotient:

$$\frac{\partial f(x)}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- ▶ Even better: Symmetric difference quotient


$$\frac{\partial f(x)}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x-h)}{2h}$$



# Numerical Differentiation

How to choose  $h$ ?

- ▶ For  $h = 0$  the expression is undefined
- ▶ Choose  $h$  to trade-off:
  - ▶ Rounding error (finite precision)
  - ▶ Approximation error (wrong)
- ▶ Good choice:  $\sqrt[3]{\epsilon}$  with  $\epsilon$  the machine precision
- ▶ Examples:
  - ▶  $\epsilon = 6 \times 10^{-8}$  for single precision (32 bit)
  - ▶  $\epsilon = 1 \times 10^{-16}$  for double precision (64 bit)





[en.wikipedia.org/wiki/  
Numerical\\_differentiation](https://en.wikipedia.org/wiki/Numerical_differentiation)

# Numerical Differentiation

Example: Sigmoid derivative using symmetric differences with single precision:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\frac{\partial \sigma(x)}{\partial x} = \sigma(x)(1 - \sigma(x))$$







# Numerical Differentiation

Example: Sigmoid derivative using symmetric differences with single precision:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



$$\frac{\partial \sigma(x)}{\partial x} = \sigma(x)(1 - \sigma(x))$$



# Numerical Differentiation

Example: Sigmoid derivative using symmetric differences with single precision:

$$\sigma(x) = \frac{1}{1 + e^{-x}} \qquad \frac{\partial \sigma(x)}{\partial x} = \sigma(x)(1 - \sigma(x))$$



## 4.3

# Preprocessing and Initialization

# Data Preprocessing

# Data Preprocessing


**Remember what happens for positive inputs:**

$$g(x) = g \left( \sum_i a_i x_i + b \right)$$


The gradient wrt. parameter  $a_i$  is given by:

$$\frac{\partial \mathcal{L}}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} x_i$$

- ▶ Both **terms in blue** are positive
- ▶ All gradients have the same sign (+ or -)
- ▶ We should pre-process the input data such that it is “well distributed”






# Data Preprocessing





- **Zero-center:**  $x_{i,j} \leftarrow x_{i,j} - \mu_j$  with  $\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$
- **Normalization:**  $x_{i,j} \leftarrow x_{i,j} / \sigma_j$  with  $\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$

# Data Preprocessing



- **Decorrelate:** Multiply with eigenvectors of covariance matrix
- **Whiten:** Divide by square root of eigenvalues of covariance matrix

# Data Preprocessing



- Classification loss becomes less sensitive to changes in the weight matrix



# Data Preprocessing

## Common Practices for Images:

- ▶ AlexNet: Subtract mean image  
(mean image:  $W \times H \times 3$  numbers)
- ▶ VGGNet: Subtract per-channel mean  
(mean along each channel: 3 numbers)
- ▶ ResNet: Subtract per-channel mean and divide by per-channel std. dev.  
(mean along each channel: 3 numbers)
- ▶ Whitening is less common

# Weight Initialization

# Recap: Stochastic Gradient Descent


**Algorithm** for training an MLP using (stochastic) gradient descent:

1. Initialize weights  $\mathbf{w}$ , pick learning rate  $\eta$  and minibatch size  $|\mathcal{X}_{\text{batch}}|$
2. Draw (random) minibatch  $\mathcal{X}_{\text{batch}} \subseteq \mathcal{X}$
3. For all elements  $(\mathbf{x}, \mathbf{y}) \in \mathcal{X}_{\text{batch}}$  of minibatch (in parallel) do:
  - 3.1 Forward propagate  $\mathbf{x}$  through network to calculate  $\mathbf{h}_1, \mathbf{h}_2, \dots, \hat{\mathbf{y}}$
  - 3.2 Backpropagate gradients through network to obtain  $\nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$
4. Update gradients:  $\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{1}{|\mathcal{X}_{\text{batch}}|} \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{X}_{\text{batch}}} \nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$
5. If validation error decreases, go to step 2, otherwise stop

**Question:**

- How to best initialize the weights  $\mathbf{w}$ ?

# Constant Initialization



- ▶ How to initialize the parameters  $\mathbf{w}$  of all network layers?
- ▶ Simple solution: set all network parameters to a constant (i.e.,  $\mathbf{w} = 0$ )
- ▶ Learning not be possible (all units of each layer are learning the same)

# Weight Initialization

Consider a layer in a Multi-Layer Perceptron:

$$g(x) = g \left( \sum_i a_i x_i + b \right)$$


The gradient wrt. parameter  $a_i$  is given by:

$$\frac{\partial \mathcal{L}}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} x_i$$

Remark:

- For  $g(\cdot)$ , we will use Tanh and ReLU in the following

# Small Random Numbers




## Tanh Activation Function:

- ▶ Draw weights independently from Gaussian with small std. dev ( $\sigma = 0.01$ )
- ▶ Activations (=activation function outputs) in deeper layers tend towards zero
- ▶ Gradients wrt. weights thus also tend towards zero  $\Rightarrow$  no learning:

$$\frac{\partial \mathcal{L}}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} x_i = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} 0 = 0$$

# Large Random Numbers




## Tanh Activation Function:

- ▶ Draw weights independently from Gaussian with large std. dev ( $\sigma = 0.2$ )
- ▶ All activation functions saturate
- ▶ Local gradients are all becoming zero  $\Rightarrow$  no learning:

$$\frac{\partial \mathcal{L}}{\partial a_i} = \frac{\partial \mathcal{L}}{\partial g} \frac{\partial g}{\partial x} x_i = \frac{\partial \mathcal{L}}{\partial g} 0 x_i = 0$$

# Xavier Initialization



## Tanh Activation Function:

- ▶ Glorot et al. draw weights independently from Gaussian with  $\sigma^2 = 1/D_{in}$
- ▶  $D_{in}$  denotes the dimension of the input to the layer, may vary across layers
- ▶ Activation distribution now well scaled across all layers



# Xavier Initialization


Why  $\sigma = 1/\sqrt{D_{in}}$ ? Let us consider  $y = g(\mathbf{w}^\top \mathbf{x})$  and assume that all  $x_i$  and  $w_i$  are independent and identically (i.i.d.) distributed with zero mean. Let further  $g'(0) = 1$ . Then:

$$\begin{aligned}\text{Var}(y) &\approx \text{Var}(\mathbf{w}^\top \mathbf{x}) = D_{in} \text{Var}(x_i w_i) \\ &= D_{in} (\mathbb{E}[x_i^2 w_i^2] - \mathbb{E}[x_i w_i]^2) \\ &= D_{in} (\mathbb{E}[x_i^2] \mathbb{E}[w_i^2] - \mathbb{E}[x_i]^2 \mathbb{E}[w_i]^2) \\ &= D_{in} \mathbb{E}[x_i^2] \mathbb{E}[w_i^2] \\ &= D_{in} \text{Var}(x_i) \text{Var}(w_i)\end{aligned}$$

Thus:

$$\text{Var}(w_i) = 1/D_{in} \Rightarrow \text{Var}(y) = \text{Var}(x_i)$$


# Xavier Initialization



## ReLU Activation Function:

- ▶ Xavier initialization assumes zero centered activation function
- ▶ For ReLU, activations again start collapsing to zero for deeper layers

# He Initialization



## ReLU Activation Function:

- ▶ Since ReLU is restricted to positive outputs, variance must be **doubled**
- ▶ He et al. draw weights independently from Gaussian with  $\sigma^2 = 2/D_{in}$
- ▶ Activation distribution now well scaled across all layers

# Summary

## **Data Preprocessing:**

- ▶ Zero-centering the network inputs is important for efficient learning
- ▶ Decorrelation and whitening used less frequently

## **Weight Initialization:**

- ▶ Proper initialization important for ensuring a good “gradient flow”
- ▶ For zero-centered activation functions, use Xavier initialization
- ▶ For ReLU activation functions, use He initialization
- ▶ Initialization is a research topic, much more literature on this topic